

COMET LEAP Token Burn And Revenue Convergence v1

Generated: 2026-05-26

Local master repo: `C:\tmp\lantern-os`

Canonical remote link to use in PDFs: `https://github.com/alex-place/lantern-os`

Status: operator-review draft, built with the Lantern OS 12-step convergence loop.

Executive Answer

The previous low-confidence answer was weak because it mixed shallow market searches, token-cost claims, and revenue projections without enough independent evidence classes. This version raises confidence only where the evidence is strong:

- **High confidence:** token unit costs, repo artifact existence, current dirty repo state, service streams that can be sold manually now.
- **Medium confidence:** 30-90 day revenue targets from founder-led service packaging.
- **Low confidence:** scaled 12-month revenue without live conversion data.

The fastest money is not a token economy. It is founder-led paid installation, caregiver/accessibility setup, homeschool/local AI tutoring setup, and COMET LEAP report packaging. Token issuance remains a later financing instrument, not a first revenue stream.

Source Base

Source	Use	Confidence Impact
OpenAI official API pricing, <code>`https://openai.com/api/pricing/`</code> and <code>`https://platform.openai.com/docs/pricing/`</code>	Current per-token costs and batch/cached pricing	High
Anthropic official Claude pricing, <code>`https://platform.claude.com/docs/en/about-claude/pricing`</code>	Claude input/output/cache/batch costs	High
Google Gemini pricing/docs, <code>`https://ai.google.dev/gemini-api/docs/pricing`</code> and <code>`https://ai.google.dev/gemini-api/docs/tokens`</code>	Gemini token billing and low-cost Flash tiers	High
Berkshire Hathaway annual letters / owner earnings concept, <code>`https://www.berkshirehathaway.com/`</code>	Owner earnings lens: cash that can be removed without harming the business	Medium-high
Pew homeschooling summary, <code>`https://www.pewresearch.org/short-reads/2025/02/20/a-look-at-homeschooling-in-the-us/`</code>	Homeschool market baseline	Medium
Johns Hopkins homeschool growth 2024-2025, <code>`https://education.jhu.edu/edpolicy/policy-research-initiatives/homeschool-hub/homeschool-growth-2024-2025/`</code>	Directional homeschool growth	Medium
AARP 2024 tech trends 50+, <code>`https://www.aarp.org/pri/topics/technology/internet-media-devices/2024-technology-trends-older-adults.html`</code>	Older-adult tech adoption and AI skepticism	Medium
AARP telehealth 50+, <code>`https://www.aarp.org/pri/topics/health/coverage-access/telehealth-health-care/`</code>	Caregiver/health-adjacent adoption surface	Medium

Method Upgrade

Old method: three-ish web searches plus projection.

New method:

1. Use official API pricing for token costs.
2. Separate **token burn** from **cash burn**.
3. Use Buffett-style owner earnings: revenue minus operating costs minus maintenance reinvestment.
4. Use persona streams only where a payment path exists now.
5. Score confidence from evidence class, not optimism.
6. Prefer manual service revenue before SaaS scaling claims.
7. Treat crypto/token issuance as optional financing, not core revenue.
8. Split past, present, and future into 12 convergence frames.
9. Record blockers and next build actions.
10. Attach repo and PDF links.
11. Run local convergence loop.
12. Promote only after operator review.

Token Burn Cost Model

Assumptions per active family per month:

- 200 short local chats.
- 40 cloud-assisted chats.
- Average cloud chat: 2,000 input tokens, 600 output tokens.
- Monthly cloud tokens per family: 80,000 input, 24,000 output.
- Local-first default sends only sensitive summaries or opt-in escalations to cloud.

Per-Family Monthly Token Cost

Provider / Model Class	Input \$/1M	Output \$/1M	Cost Per Family / Month	Notes
OpenAI nano-class	0.05	0.40	\$0.0136	Best for classification, routing, simple summaries
OpenAI mini-class	0.25	2.00	\$0.0680	Good default cloud assist budget
OpenAI 4.1 mini-class	0.40	1.60	\$0.0704	Strong low-cost fallback
Anthropic Haiku-class	1.00	5.00	\$0.2000	Higher quality small-model path
Anthropic Sonnet-class	3.00	15.00	\$0.6000	Use only for high-stakes synthesis/review
Gemini Flash-class batch	0.25	1.50	\$0.0560	Strong batch/back-office candidate

Token Burn By Customer Count

Families	Lite Cost / Mo	Balanced Cost / Mo	Premium Review Cost / Mo
1	\$0.01	\$0.07	\$0.60
10	\$0.14	\$0.68	\$6.00
50	\$0.68	\$3.40	\$30.00
100	\$1.36	\$6.80	\$60.00
500	\$6.80	\$34.00	\$300.00
1,000	\$13.60	\$68.00	\$600.00

Conclusion: **token burn is not the near-term threat. Human support time, installation friction, trust, and distribution are the burn risks.**

Buffett-Style Owner Earnings Lens

Owner earnings proxy:

Owner earnings =
 cash collected
 - direct API/hosting/tool costs
 - customer support labor
 - required maintenance reinvestment

For Lantern OS, the best early owner-earnings streams are high-trust services with low token cost:

- installation/setup;
- accessibility/caregiver configuration;
- homeschool/private learning setup;
- COMET LEAP report packaging;
- local RAG/IP compression.

12 Convergence Confidence Tables

Confidence scale:

- 90-100: directly verified local or official source.
- 75-89: strong evidence and immediate build path.
- 60-74: plausible, needs live conversion data.
- 40-59: speculative scaling case.
- below 40: idea only.

1. Past Artifact Proof

Change / Stream	Evidence	Confidence	Action
COMET LEAP 30-day artifacts exist	Source repo artifact manifest	95	Promote only selected files
30 PNG set exists	`day_01.png` through `day_30.png` observed earlier	95	Add checksum manifest
PDFs validate as PDF	`%PDF-` header checks performed earlier	95	Regenerate master PDF with repo link

2. Present Repo State

Change / Stream	Evidence	Confidence	Action
Lantern OS repo exists	`C:\tmp\lantern-os` git repo	99	Continue here
Source repos are dirty	Local git status showed HFF and orchestrator changes	95	Do not reset or sync
Remote URL not configured locally	`git remote -v` returned no remote	95	Add link in PDF, set remote only on approval

3. Windows Surface Revenue

Offer	Price	Build Now?	Confidence
Lantern Windows setup bundle	\$49-\$149 one-time	Yes	85
Accessibility desktop setup	\$99-\$299 one-time	Yes	82
Family onboarding call	\$50-\$150	Yes	80

4. Dual Boot / NixOS Revenue

Offer	Price	Build Now?	Confidence
Dual boot planning packet	\$49-\$99	Yes	78
Guided install session	\$199-\$499	Operator physical action required	65
Managed NixOS config review	\$99-\$299	Yes	72

5. Homeschool / Education Stream

Offer	Price	Build Now?	Confidence
Local-first AI tutor setup	\$25-\$75/mo	Yes, manual	74
Homeschool co-op package	\$150-\$500/mo	Needs pilot	62
Printable learning packet	\$19-\$49	Yes	76

6. Caregiver / Accessibility Stream

Offer	Price	Build Now?	Confidence
Caregiver tech simplification	\$99-\$299	Yes	82
Voice/local chat setup	\$25-\$75/mo	Yes	70
Telehealth prep assistant	\$49-\$149	Yes, non-medical boundary required	68

7. COMET LEAP Reports Stream

Offer	Price	Build Now?	Confidence
Founder confidence report	\$49-\$199	Yes	80
Cash runway report	\$99-\$299	Yes	78
30-day convergence packet	\$199-\$999	Yes, review quality first	72

8. Local RAG / IP Compression Stream

Offer	Price	Build Now?	Confidence
IP corpus compression	\$199-\$999	Yes	78
Spin-state mini RAG	\$99-\$499	Yes	75
Repo-to-claim-index conversion	\$299-\$1,500	Yes	76

9. Privacy AI Assistant Stream

Offer	Price	Build Now?	Confidence
Privacy-first assistant install	\$49-\$199	Yes	74
Local model routing setup	\$99-\$499	Yes	70
Small-org private assistant	\$500-\$2,500	Needs trust + support	58

10. Subscription Stream

Offer	Price	Build Now?	Confidence
Family Lite	\$9-\$19/mo	Needs payment surface	65
Family Plus	\$25-\$49/mo	Needs support loop	62
Co-op / group	\$99-\$299/mo	Needs pilot	55

11. 30/60/90-Day Cash Path

Window	Target Revenue	Confidence	Key Constraint
30 days	\$250-\$1,000	75	Manual selling and setup
60 days	\$1,000-\$3,000	66	Repeatable offer page and referrals
90 days	\$3,000-\$7,500	55	Support capacity and proof

12. Future Scale Path

Horizon	Target Revenue	Confidence	Gate
6 months	\$5k-\$15k MRR	45	Productized onboarding
12 months	\$15k-\$50k MRR	35	Retention and distribution data
24 months	\$50k+ MRR	25	Team, channels, compliance posture

Money Streams To Build Now

Build these first because they need little new code and have direct payment logic:

- 1. ****Windows Lantern Surface Setup****: package existing icon, shortcuts, docs, and local launchers. Sell as setup/help.
- 2. ****COMET LEAP Founder Report Pack****: cash runway, confidence report, 30-day model, truth-only report.
- 3. ****RAG/IP Compression Sprint****: convert messy repo/docs into claim-indexed mini states.

4. **Caregiver/Accessibility Setup**: local-first simplification with explicit non-medical boundary.
5. **Dual Boot Planning Packet**: paid planning and checklist; physical install remains operator-controlled.

Token Strategy

Do not launch a crypto token now.

If "token" means LLM token burn:

- use local-first routing;
- batch non-urgent synthesis;
- cache system prompts;
- use nano/mini models for routing;
- reserve Sonnet/frontier models for review and evidence synthesis.

If "token" means crypto/tokenomics:

- hold until there is revenue, a community, and legal review;
- avoid promises of yield, profit, governance authority, or investment return;
- use credits or prepaid service units before any public token instrument.

One-Shot Decision

The higher-confidence build path is:

Week 1: paid setup + report pack
Week 2: RAG/IP compression sprint
Week 3: caregiver/homeschool pilot
Week 4: subscription checkout only if manual buyers exist

This keeps owner earnings positive, avoids token theater, and turns the existing COMET LEAP artifacts into cashable services before scaling software.